

PetClassification_CNN

June 23, 2023

0.1 Pet classification using CNN

```
[ ]: # To mount the drive
from google.colab import drive
drive.mount('/content/gdrive')
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).

```
[ ]: import tensorflow as tf
from keras.models import Sequential
from keras.layers import Conv2D
from keras.layers import MaxPooling2D
from keras.layers import Flatten
from keras.layers import Dense, Dropout, BatchNormalization
from keras.preprocessing.image import ImageDataGenerator
```

```
[ ]: # To generate more images

from keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(rotation_range=13,
                                   width_shift_range=0.3,
                                   height_shift_range=0.3,
                                   shear_range=0.2,
                                   zoom_range=0.1,
                                   rescale=1./255,
                                   horizontal_flip=True,
                                   vertical_flip=True,
                                   validation_split=0.20,
                                   fill_mode='nearest')
```

```
[ ]: test_datagen = ImageDataGenerator(rescale = 1./255)
```

```
[ ]: training_set = train_datagen.flow_from_directory('/content/gdrive/MyDrive/Pets/
↳train',
                                                    target_size = (128,128),
```

```
batch_size = 8,  
class_mode = 'categorical')
```

Found 40 images belonging to 2 classes.

```
[ ]: test_set = test_datagen.flow_from_directory('/content/gdrive/MyDrive/Pets/test',  
                                              target_size = (128,128),  
                                              batch_size = 8,  
                                              class_mode = 'categorical')
```

Found 20 images belonging to 2 classes.

```
[ ]: model=Sequential()  
model.add(Conv2D( 32, (5, 5 ), activation = 'relu', padding = 'same',  
    ↪input_shape = (128, 128 , 3 ),kernel_regularizer =tf.keras.regularizers.l2(  
    ↪l=0.03)))  
# model.add(BatchNormalization())  
model.add(MaxPooling2D(2,2))  
model.add(Conv2D(64, (5, 5 ), activation = 'relu',kernel_regularizer =tf.keras.  
    ↪regularizers.l2( l=0.03)))  
# model.add(BatchNormalization())  
model.add(MaxPooling2D((2,2)))  
model.add(Dropout(0.4))  
model.add(Flatten())  
model.add(Dense(32, activation = 'relu',kernel_regularizer =tf.keras.  
    ↪regularizers.l2( l=0.001)))  
model.add(Dense(2, activation = 'softmax'))
```

```
[ ]: model.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics =  
    ↪['accuracy'])
```

Running for 100 epochs

```
[ ]: ! pip install livelossplot
```


Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>
Requirement already satisfied: livelossplot in /usr/local/lib/python3.10/dist-packages (0.5.5)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from livelossplot) (3.7.1)
Requirement already satisfied: bokeh in /usr/local/lib/python3.10/dist-packages (from livelossplot) (2.4.3)
Requirement already satisfied: Jinja2>=2.9 in /usr/local/lib/python3.10/dist-packages (from bokeh->livelossplot) (3.1.2)
Requirement already satisfied: numpy>=1.11.3 in /usr/local/lib/python3.10/dist-packages (from bokeh->livelossplot) (1.22.4)
Requirement already satisfied: packaging>=16.8 in

/usr/local/lib/python3.10/dist-packages (from bokeh->livelossplot) (23.1)
 Requirement already satisfied: pillow>=7.1.0 in /usr/local/lib/python3.10/dist-packages (from bokeh->livelossplot) (8.4.0)
 Requirement already satisfied: PyYAML>=3.10 in /usr/local/lib/python3.10/dist-packages (from bokeh->livelossplot) (6.0)
 Requirement already satisfied: tornado>=5.1 in /usr/local/lib/python3.10/dist-packages (from bokeh->livelossplot) (6.3.1)
 Requirement already satisfied: typing-extensions>=3.10.0 in /usr/local/lib/python3.10/dist-packages (from bokeh->livelossplot) (4.5.0)
 Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->livelossplot) (1.0.7)
 Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->livelossplot) (0.11.0)
 Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->livelossplot) (4.39.3)
 Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->livelossplot) (1.4.4)
 Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->livelossplot) (3.0.9)
 Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->livelossplot) (2.8.2)
 Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2>=2.9->bokeh->livelossplot) (2.1.2)
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib->livelossplot) (1.16.0)

```
[ ]: from livelossplot import PlotLossesKerasTF
```

```
[ ]: tf.config.run_functions_eagerly(True)
```

```
[ ]: pets=model.fit(training_set,epochs=100,validation_data =_
    ↪test_set,callbacks=[PlotLossesKerasTF()])
```



```

accuracy
      training      (min: 0.425, max: 0.775, cur: 0.700)
    validation      (min: 0.350, max: 0.750, cur: 0.650)
Loss
      training      (min: 0.584, max: 2.154, cur: 0.584)
    validation      (min: 0.630, max: 1.894, cur: 0.636)
5/5 [=====] - 1s 235ms/step - loss: 0.5843 - accuracy:
0.7000 - val_loss: 0.6363 - val_accuracy: 0.6500

```

Doing hyperparameter tuning and adding l2 regularization

```
[ ]: test_loss, test_accuracy = model.evaluate(test_set)
```

```

2/2 [=====] - 0s 24ms/step - loss: 2.8962 - accuracy:
0.5500

```


```
[ ]: model1=Sequential()
model1.add(Conv2D( 32, (5, 5 ), activation = 'relu', padding = 'same',
↳input_shape = (128, 128 , 3 ),kernel_regularizer =tf.keras.regularizers.l1(
↳l=0.01)))
# model.add(BatchNormalization())
model1.add(MaxPooling2D(2,2))
model1.add(Conv2D(64, (5, 5 ), activation = 'relu',kernel_regularizer =tf.keras.
↳regularizers.l1( l=0.01)))

```

```
# model.add(BatchNormalization())
model1.add(MaxPooling2D((2,2)))
model1.add(Dropout(0.4))
model1.add(Flatten())
model1.add(Dense(32, activation = 'relu',kernel_regularizer =tf.keras.
↳regularizers.l1( l=0.01)))
model1.add(Dense(2, activation = 'softmax'))
```

```
[ ]: model1.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics =_
↳['accuracy'])
```

```
[ ]: pets=model1.fit(training_set,epochs=100,validation_data =_
↳test_set,callbacks=[PlotLossesKerasTF()])
```



```
accuracy
      training      (min:    0.425, max:    1.000, cur:    1.000)
    validation      (min:    0.450, max:    0.750, cur:    0.600)

Loss
      training      (min:    0.231, max:    3.992, cur:    0.233)
    validation      (min:    1.073, max:    3.415, cur:    3.290)
3/3 [=====] - 2s 885ms/step - loss: 0.2334 - accuracy:
1.0000 - val_loss: 3.2896 - val_accuracy: 0.6000
```

```
[ ]: test_loss, test_accuracy = model1.evaluate(test_set)
```


```
2/2 [=====] - 0s 31ms/step - loss: 3.2896 - accuracy: 0.6000
```

1 Running for 200 epochs

```
[ ]: model2=Sequential()  
model2.add(Conv2D( 32, (5, 5 ), activation = 'relu', padding = 'same',  
    ↳input_shape = (256, 256 , 3 )))  
model2.add(MaxPooling2D(2,2))  
model2.add(Conv2D(64, (5, 5 ), activation = 'relu',kernel_regularizer =tf.keras.  
    ↳regularizers.l2( l=0.01)))  
# model2.add(BatchNormalization())  
model2.add(MaxPooling2D((2,2)))  
model2.add(Dropout(0.4))  
model2.add(Flatten())  
model2.add(Dense(32, activation = 'relu',kernel_regularizer =tf.keras.  
    ↳regularizers.l2( l=0.01)))  
# model2.add(BatchNormalization())  
model2.add(Dense(2, activation = 'softmax'))
```

```
[ ]: model2.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics =  
    ↳['accuracy'])
```

```
[ ]: pets=model2.fit(training_set,epochs=200,validation_data =  
    ↳test_set,callbacks=[PlotLossesKerasTF()])
```



```

accuracy
      training      (min: 0.500, max: 1.000, cur: 0.975)
    validation      (min: 0.350, max: 0.750, cur: 0.450)
Loss
      training      (min: 0.242, max: 11.443, cur: 0.318)
    validation      (min: 1.213, max: 3.614, cur: 3.088)
3/3 [=====] - 2s 789ms/step - loss: 0.3180 - accuracy:
0.9750 - val_loss: 3.0880 - val_accuracy: 0.4500

```

```
[ ]: test_loss, test_accuracy = model2.evaluate(test_set)
```

```

2/2 [=====] - 0s 25ms/step - loss: 3.0880 - accuracy:
0.4500

```

2 Running for 300 epochs

trying without regularization and using SGD optimizer

```

[ ]: model3=Sequential()
    model3.add(Conv2D( 32, (5, 5 ), activation = 'relu', padding = 'same',
        ↪input_shape = (256, 256 , 3 )))
    model3.add(MaxPooling2D(2,2))

```

```

model3.add(Conv2D(64, (5, 5), activation = 'relu',kernel_regularizer =tf.keras.
↳regularizers.l2( l=0.01)))
# model2.add(BatchNormalization())
model3.add(MaxPooling2D((2,2)))
model3.add(Dropout(0.4))
model3.add(Flatten())
model3.add(Dense(32, activation = 'relu',kernel_regularizer =tf.keras.
↳regularizers.l2( l=0.01)))
# model2.add(BatchNormalization())
model3.add(Dense(2, activation = 'softmax'))

```

```

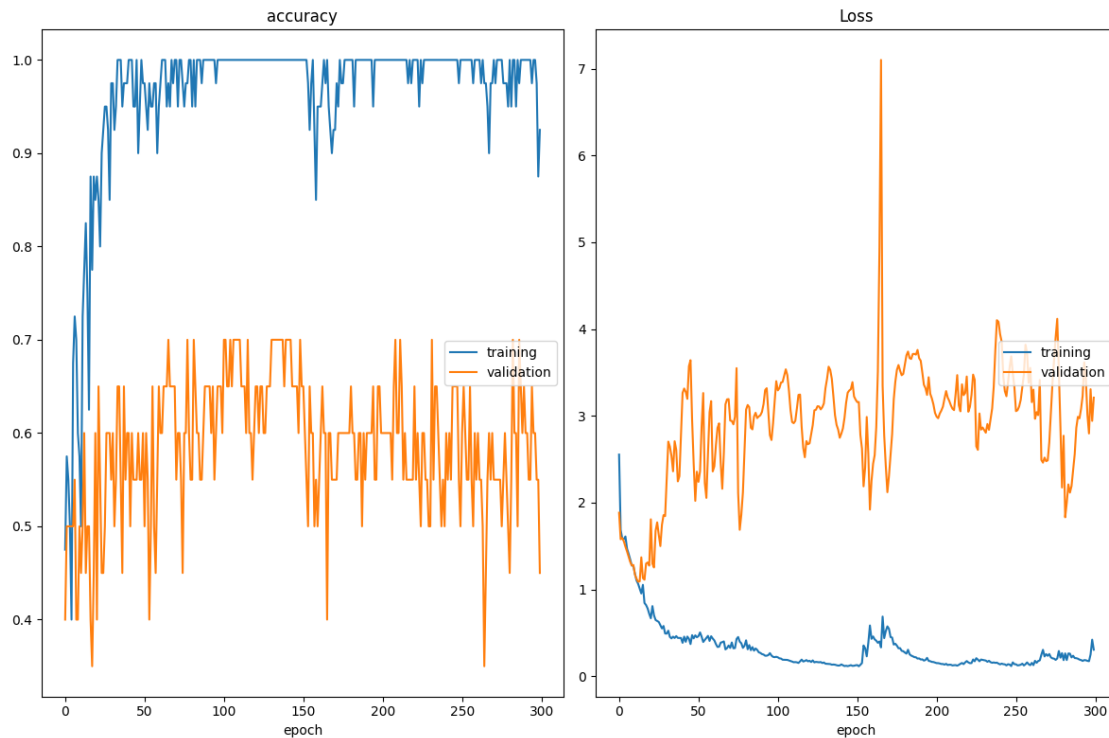
[ ]: model3.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics =_
↳['accuracy'])

```

```

[ ]: pets=model3.fit(training_set,epochs=300,validation_data =_
↳test_set,callbacks=[PlotLossesKerasTF()])

```



accuracy

training	(min: 0.400, max: 1.000, cur: 0.925)
validation	(min: 0.350, max: 0.700, cur: 0.450)


Loss

training	(min: 0.115, max: 2.554, cur: 0.306)
validation	(min: 1.090, max: 7.101, cur: 3.210)

3/3 [=====] - 1s 531ms/step - loss: 0.3060 - accuracy: 0.9250 - val_loss: 3.2102 - val_accuracy: 0.4500

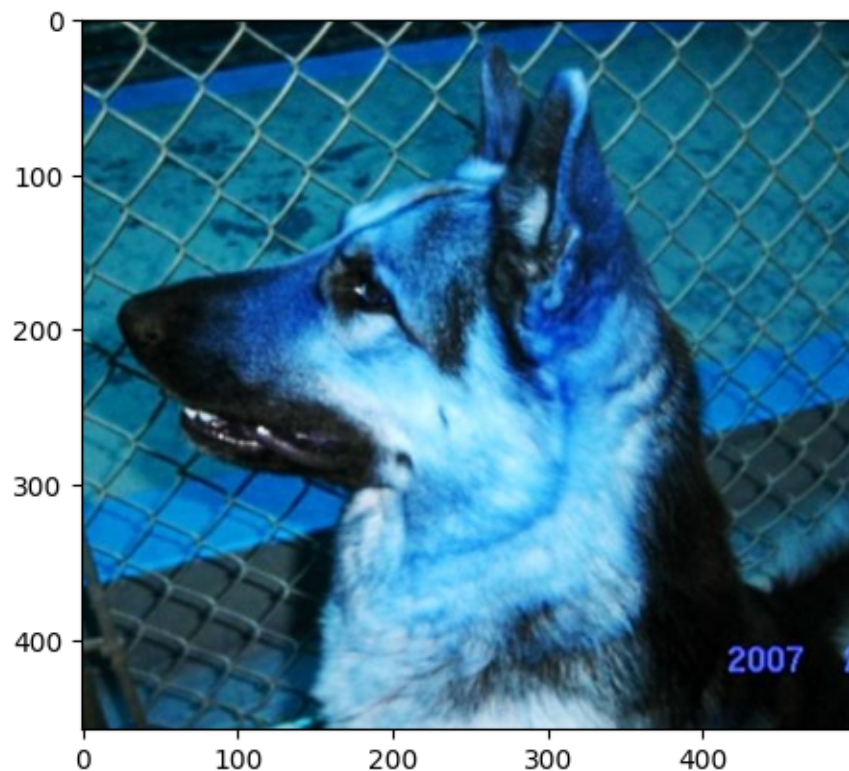
```
[ ]: import cv2
import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline
```

```
[ ]: image = '/content/gdrive/MyDrive/Pets/test/cats/110.jpg'
img = cv2.imread(image)
plt.imshow(img)
plt.show()
img = cv2.resize(img,(256,256))
img = np.reshape(img, (1, 256,256, 3))
# classes = model3.predict_classes(img)
predict_x=model3.predict(img)
classes_x=np.argmax(predict_x,axis=1)
classes=["cat" if classes_x==0 else "dog"]
print(classes)
```



1/1 [=====] - 0s 137ms/step
['cat']

```
[ ]: image = '/content/gdrive/MyDrive/Pets/test/dogs/107.jpg'
img = cv2.imread(image)
plt.imshow(img)
plt.show()
img = cv2.resize(img,(256,256))
img = np.reshape(img, (1, 256,256, 3))
# classes = model3.predict_classes(img)
predict_x=model3.predict(img)
classes_x=np.argmax(predict_x,axis=1)
classes=["dog" if classes_x==1 else "cat"]
print(classes)
```



```
1/1 [=====] - 0s 95ms/step
['dog']
```

```
[ ]:
```